



Computer Networks

Wenzhong Li, Chen Tian

Nanjing University

Material with thanks to James F. Kurose, Mosharaf Chowdhury, and other colleagues.



Outline

- TCP flow control
- TCP congestion control



TCP Flow Control



TCP header

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			



Recap: Sliding window

- Both sender and receiver maintain a **window**
- **Left edge** of window:
 - Sender: beginning of **unacknowledged** data
 - Receiver: beginning of **expected** data
 - First “hole” in received data
 - When sender gets ack, knows that receiver’s window has moved
- **Right edge**: $\text{Left edge} + \text{constant}$
 - The constant is only limited by buffer size in the transport layer

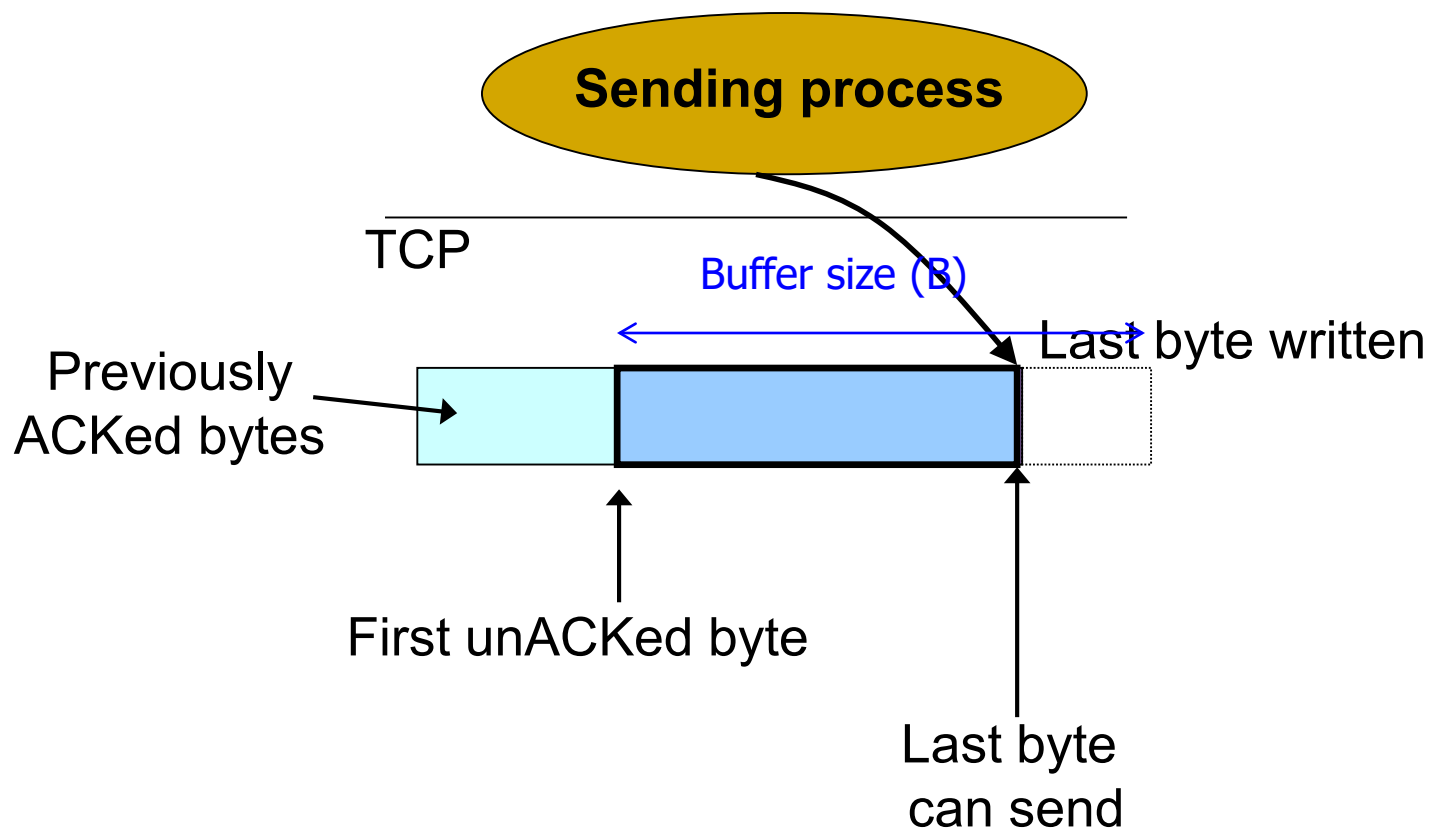


TCP: Sliding window (so far)

- Both sender and receiver maintain a **window**
- **Left edge** of window:
 - Sender: beginning of **unacknowledged** data
 - Receiver: beginning of **undelivered** data
- **Right edge**: Left edge + constant
 - The constant is only limited by buffer size in the transport layer

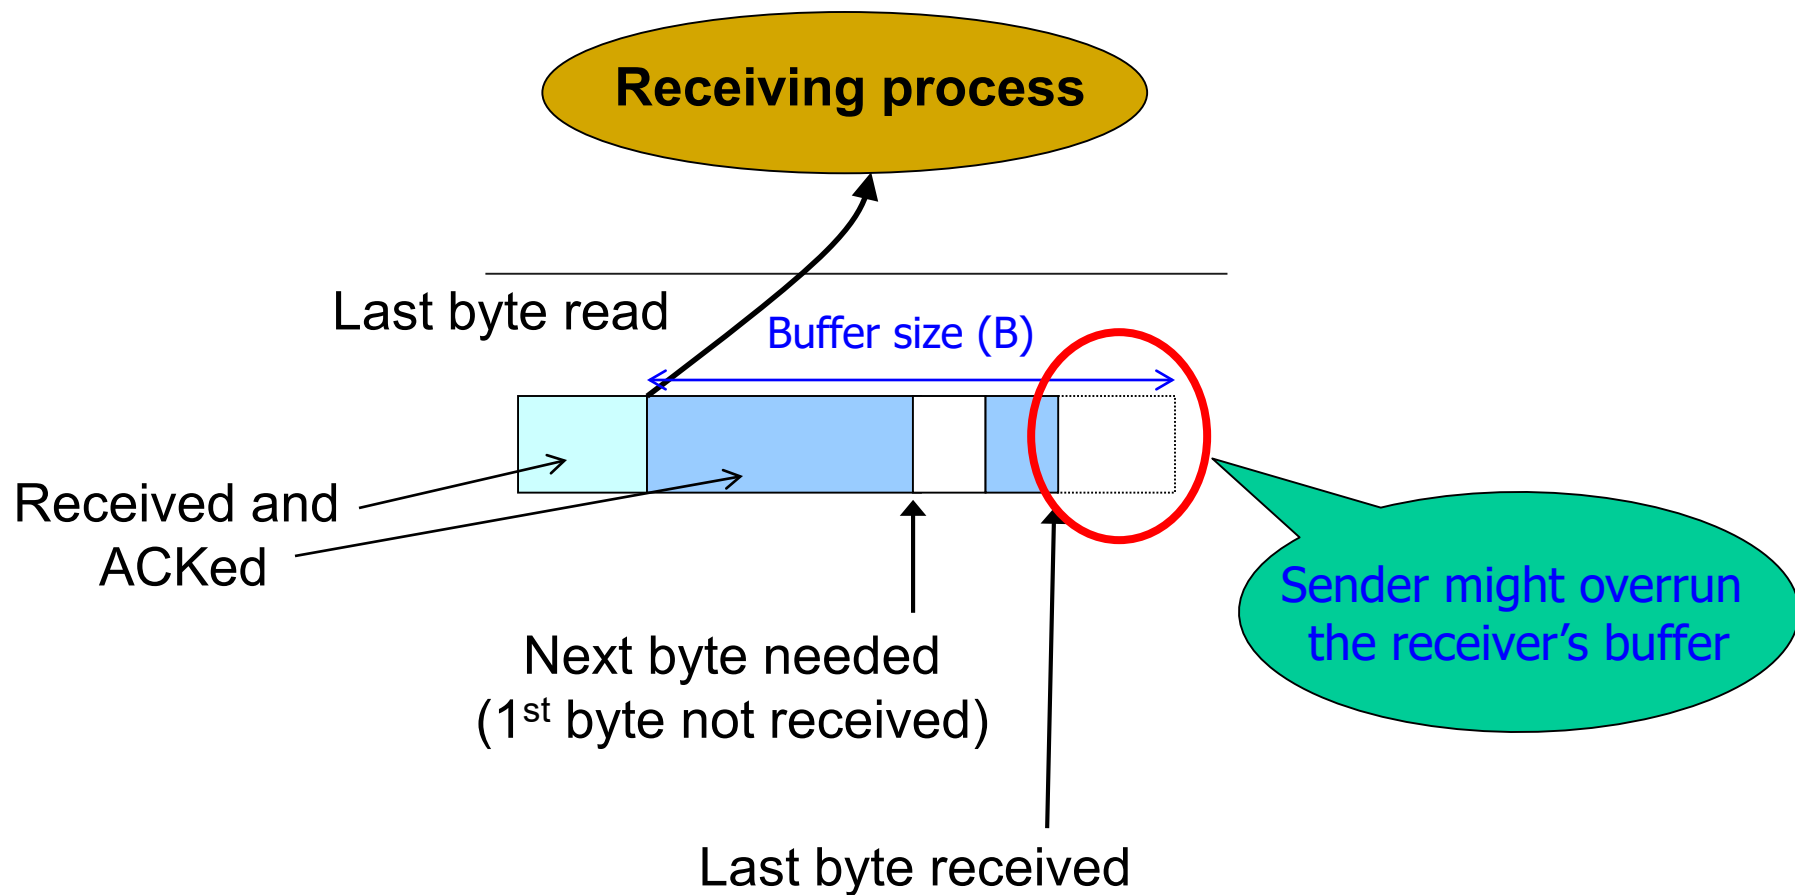


Sliding window at sender





Sliding window at receiver





Fixed sliding window?

- Fixed sliding window
 - Works well on reliable direct links
- Problem:
 - Failure to receive ACK is taken as flow control indication
 - The receiver can achieve flow control by stop sending ACK, but the sender can not distinguish between lost segment and flow control




Solution: Credit Scheme

- Receiver advertises spare room (**credits**) using an “Advertised Window” (**RWND**) to prevent sender from overflowing its window
 - Receiver indicates value of RWND in **ACKs**
 - Sender ensures that the total **number of bytes in flight** \leq **RWND**

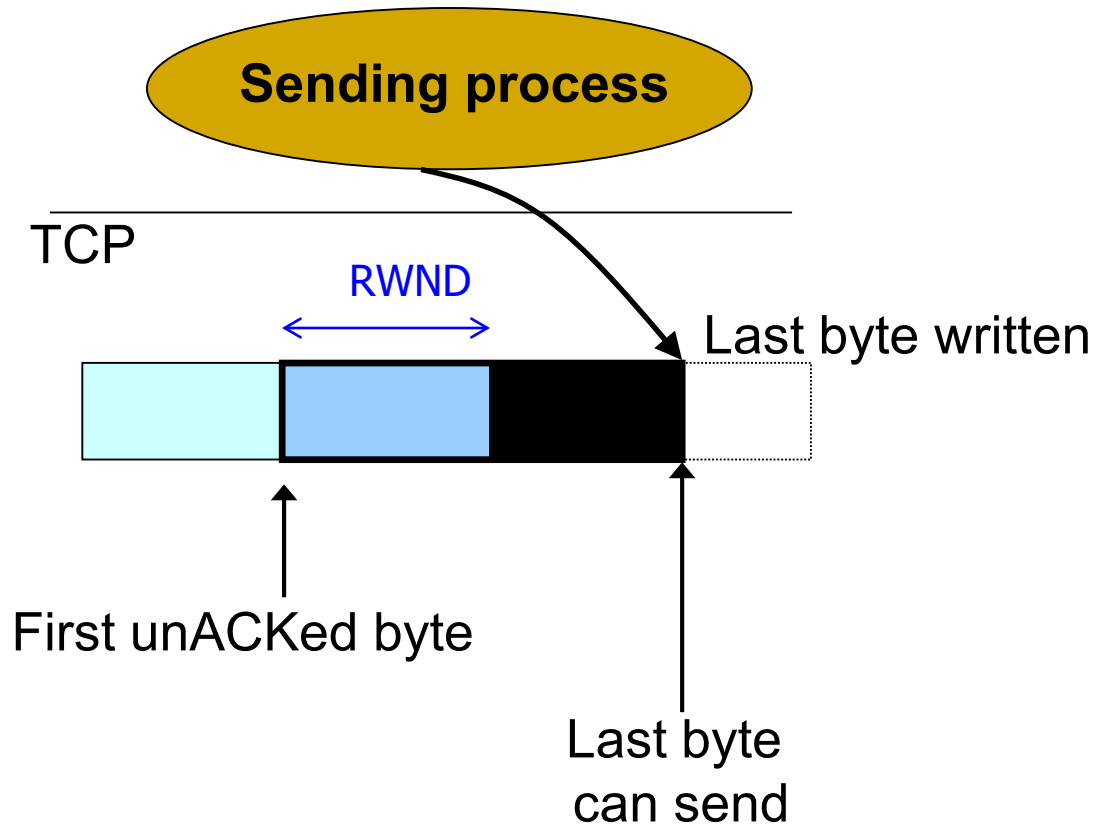


Sliding window at receiver





Sliding window at sender





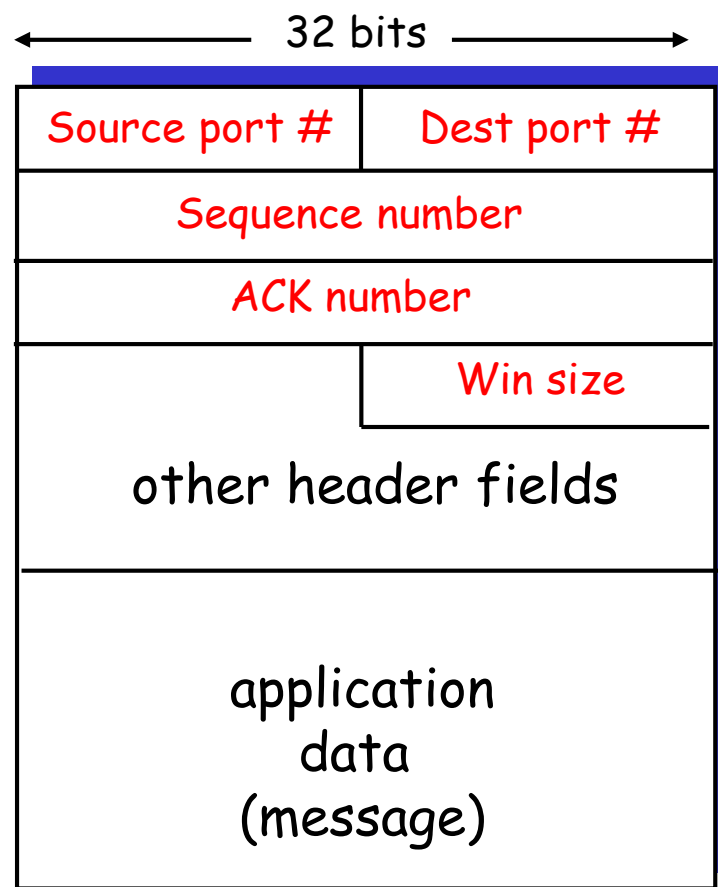
Sliding window with flow control

- **Sender**: window advances when new data ACK'd
- **Receiver**: window advances as receiving process consumes data
- Receiver advertises to the sender where the receiver window currently ends (“righthand edge”)
 - Sender agrees not to exceed this amount
- **UDP does not have flow control**
 - Data can be lost due to buffer overflow



Benefit of Credit Scheme

- Greater control on Internet
- Decouples flow control from *ACK*
 - May *ACK* without granting credit
- Each **octet** has a sequence number
- Each transport segment has seq number, ack number and window size in header



TCP segment format

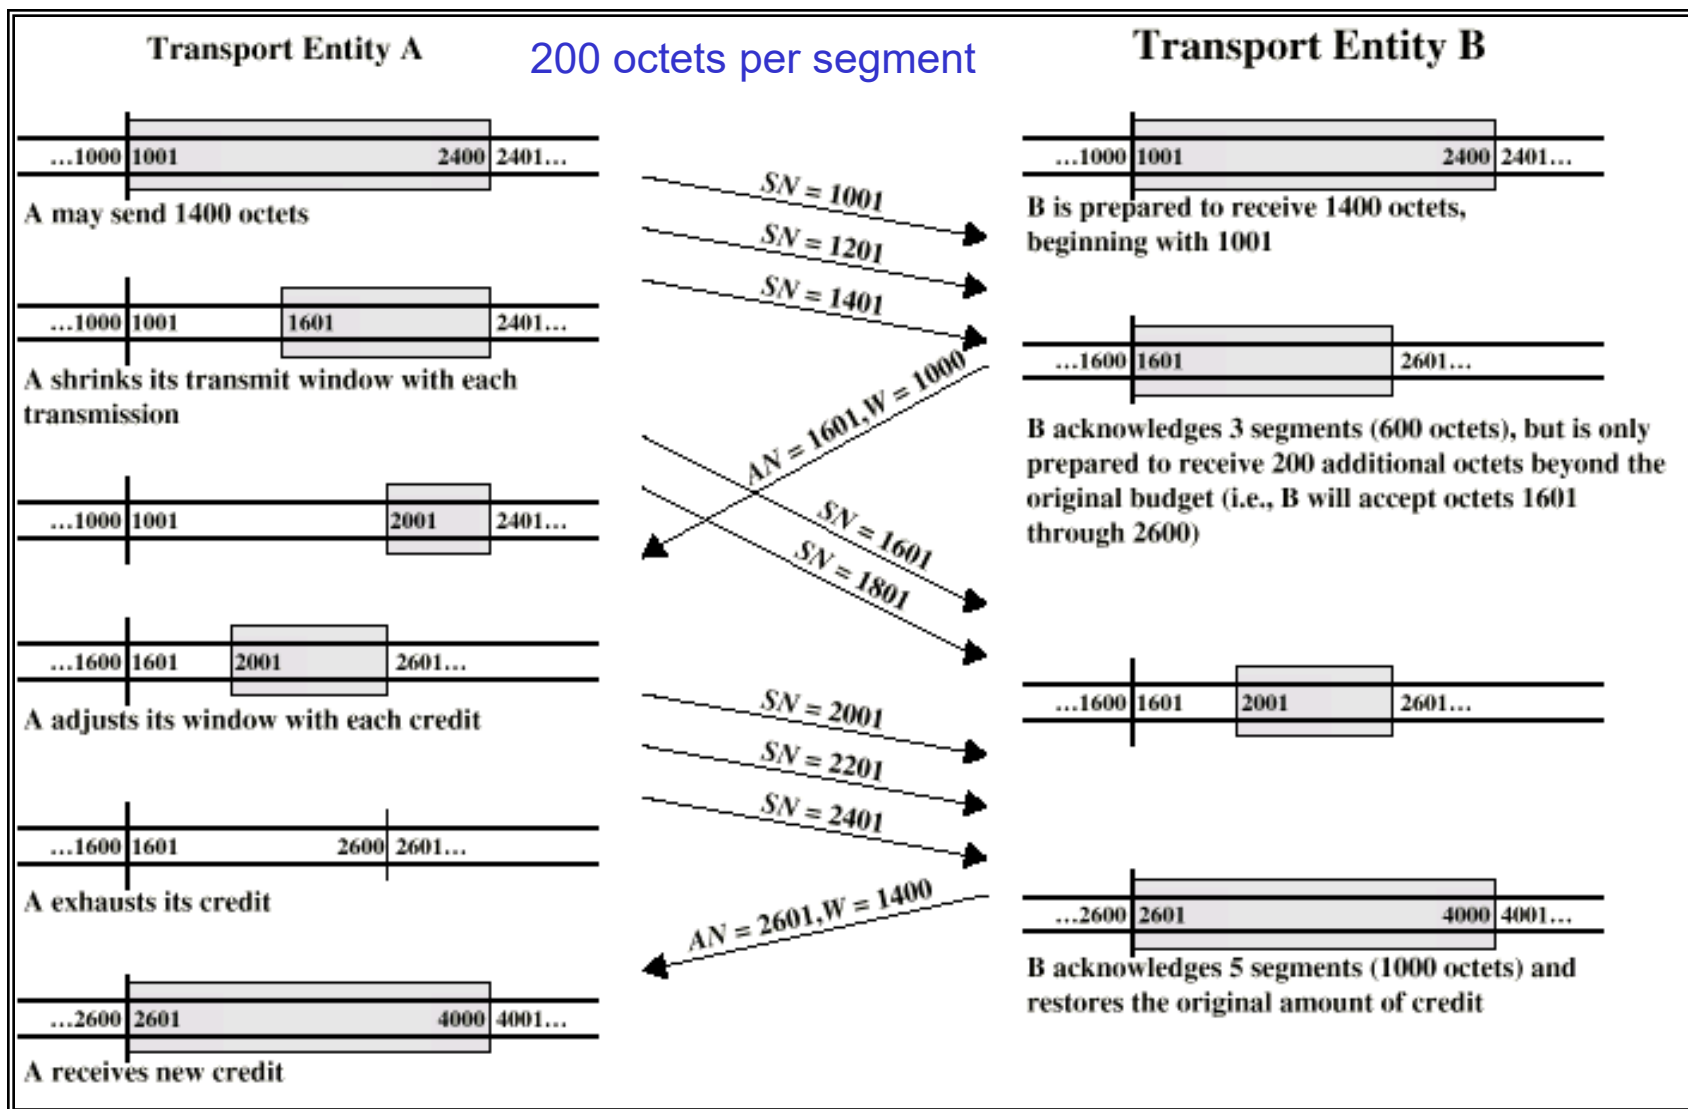


Use of Header Fields

- When sending a segment
 - seq number (SN) is that of first octet in segment
 - ACK includes $AN=i$, $W=j$
- All octets through $SN=i-1$ acknowledged
 - Next expected octet is i
- Permission to send additional window of $W=j$ octets
 - i.e. octets from i to $i+j-1$



Credit Allocation Procedure





Credit allocation flow control

Credit allocation flow control mechanism

- Suppose that the last octet of data received by B was octet number $i-1$, and that the last segment issued by B was $(AN=i, W=j)$. Then
 - To **increase credit** to an amount k ($k>j$) when no additional data have arrived, B issues $(AN=i, W=k)$
 - To acknowledge an incoming segment containing m octets of data ($m<j$) **without granting additional credit**, B issues $(AN=i+m, W=j-m)$
- If an ACK/CREDIT segment is **lost**, little harm is done. Future acknowledgments will resynchronize the protocol.
- Further, if the sender times out and retransmits a data segment, it triggers a new acknowledgment.



■ Credit allocation deadlock

- B sends A: segment with $AN=i$, $W=0$ closing rcv-window
- B sends A: $AN=i$, $W=j$ to reopen, but this maybe lost
- Now A thinks window is closed, B thinks it is open and wait

■ Handle

- Use window timer
- If timer expires without any receiving, send something
- Could be re-transmission of previous segment



TCP Congestion Control



A story: Congestion collapse in 1980s

- In 1981, TCP was standardized and widely deployed.
- No congestion control is considered.
- In October of 1986, the Internet had the first of what became a series of 'congestion collapses'. During this period, the data throughput from LBL to UC Berkeley (sites separated by 400 yards and two IMP hops) dropped from 32 Kbps to 40 bps. (Van Jacobson, Congestion Avoidance and Control)
- This open a new area of congestion control study.



Jacobson's fix to TCP

- Extend TCP's existing window-based protocol but **adapt** the window size in response to congestion
- A pragmatic and effective solution
 - Required no upgrades to routers or applications!
 - Patch of a few lines of code to TCP implementations
- Extensively researched and improved upon
 - Especially now with datacenters and cloud services



Key design considerations

- How do we know the network is congested?
 - Implicit and/or explicit signals from the network
- Who takes care of congestion?
 - End hosts (may receive some help from the network)
- How do we handle congestion?
 - Continuous adaptation



Three issues to consider

- Discovering the available (bottleneck) bandwidth
- Adjusting to variations in bandwidth
- Sharing bandwidth between flows



Abstract view



- Ignore internal structure of router and model it as a single queue for a particular input-output pair




Discovering available bandwidth



- Pick sending rate to match bottleneck bandwidth
 - Without any a priori knowledge
 - Could be gigabit link, could be a modem



Adjusting to variations in bandwidth




- Adjust rate to match instantaneous bandwidth
 - Assuming you have rough idea of bandwidth




Multiple flows and sharing bandwidth

- Two Issues:
 - Adjust total sending rate to match bandwidth
 - Allocation of bandwidth between flows





Reality



Congestion control is a resource allocation problem involving many flows, many links, and complicated global dynamics



Possible approaches

(0) Send without care

- Many packet drops



Possible approaches

(0) Send without care

(1) Reservations

- Pre-arrange bandwidth allocations
- Requires negotiation before sending packets
- Low utilization



Possible approaches

(0) Send without care

(1) Reservations

(2) Pricing

- Don't drop packets for the high-bidders
- Requires payment model



Possible approaches

- (0) Send without care
- (1) Reservations
- (2) Pricing
- (3) Dynamic Adjustment
 - Hosts **infer** level of congestion; **adjust**
 - Network **reports** congestion level to hosts; hosts **adjust**
 - Combinations of the above
 - Simple to implement but suboptimal, messy dynamics



Possible approaches

- (0) Send without care
 - (1) Reservations
 - (2) Pricing
 - (3) Dynamic Adjustment
-
- **Generality** of dynamic adjustment has proven to be very powerful
 - Doesn't presume business model, traffic characteristics, application requirements
 - But does assume good citizenship!



TCP's approach in a nutshell

- Each TCP connection has a window
 - Controls number of packets in flight
- Sending rate $\sim \text{Window}/\text{RTT}$
- Vary window size to control sending rate



Windows to keep in mind

- Congestion Window: **CWND**
 - Bytes that can be sent without overflowing routers
 - Computed by sender using congestion control algo.
- Flow control window: **RWND**
 - Bytes that can be sent without overflowing receiver
 - Determined by the receiver and reported to the sender
- Sender-side window = **$\min \{CWND, RWND\}$**
 - Assume for this lecture that $RWND \gg CWND$



Note

- This lecture talks about CWND in units of MSS
 - MSS (Maximum Segment Size): the amount of payload data in a TCP packet
 - This is only for the simplicity of presentation
- Real implementations maintain CWND in bytes



Two basic questions

- How does the sender detect congestion?
- How does the sender adjust its sending rate?
 - To address three issues
 - Finding available bottleneck bandwidth
 - Adjusting to bandwidth variations
 - Sharing bandwidth



Detecting congestion

- Packet delays
 - Tricky: noisy signal (delay often varies considerably)
- Routers tell end hosts when they're congested
- Packet loss
 - Fail-safe signal that TCP already has to detect
 - Complication: non-congestive loss (e.g., checksum errors)



Not all losses are the same

- Duplicate ACKs: isolated loss
 - Still getting ACKs
- Timeout: much more serious
 - Not enough dupacks
 - Must have suffered several losses
- Will adjust rate differently for each case



Rate adjustment

- Basic structure
 - Upon receipt of ACK (of new data): **increase rate**
 - Upon detection of loss: **decrease rate**
- How we increase/decrease the rate depends on the phase of congestion control we're in:
 - Discovering available bottleneck bandwidth
(Slow Start)
 - Adjusting to bandwidth variations
(Congestion Avoidance: AIMD)



Bandwidth discovery with “Slow Start”

- Goal: estimate available bandwidth
 - Start slow (for **safety**)
 - Ramp up quickly (for **efficiency**)
- Consider
 - $RTT = 100\text{ms}$, $MSS = 1000\text{bytes}$
 - Window size to fill 1Mbps of BW = 12.5 packets
 - Window size to fill 1Gbps = 12,500 packets
 - Either is possible!



Slow Start phase

- Sender starts at a slow rate, but **increases exponentially** until first loss
- Start with a small congestion window
 - Initially, $CWND = 1$
 - So, initial sending rate is MSS/RTT
- Double the $CWND$ for each RTT with no loss



Slow Start in action


- For each RTT: double CWND
 - i.e., for each ACK, $\text{CWND} += 1$

Linear increase per ACK ($\text{CWND}+1$) →
exponential increase per RTT ($2 * \text{CWND}$)



Slow Start in action

- For each RTT: double CWND
 - i.e., for each ACK, $CWND += 1$





When does Slow Start stop?

- Slow Start gives an estimate of available bandwidth
 - At some point, there will be loss
- Introduce a “slow start threshold” (**ssthresh**)
 - Initialized to a large value
- If $CWND > ssthresh$, stop Slow Start



Adjusting to varying bandwidth

- $CWND > ssthresh$
 - Stop rapid growth and focus on maintenance
- Now, want to track variations in this available bandwidth, oscillating around its current value
 - Repeated probing (rate increase) and backoff (decrease)
- TCP uses: “Additive Increase Multiplicative Decrease” (AIMD)



AIMD

- Additive increase: when $CWND > ssthresh$
 - For each ACK, $CWND = CWND + 1/CWND$
 - $CWND$ is increased by one only if all segments in a $CWND$ have been acknowledged
- Multiplicative decrease
- On 3 duplicate ACKs (packet loss event)
 - $ssthresh = CWND/2$
 - $CWND = ssthresh$
 - Enter Congestion Avoidance: $cwnd$ increases by 1 (linearly instead of exponentially) after each RTT
- On timeout event
 - $ssthresh = CWND/2$
 - $CWND = 1$
 - Initiate Slow Start



Illustration of Window





Figure 17.14 Illustration of Slow Start and Congestion Avoidance



Leads to the TCP “Sawtooth”



AIMD saw tooth behavior:
probing for bandwidth



Why AIMD?

- Recall the three issues
 - Finding available bottleneck bandwidth
 - Adjusting to bandwidth variations
 - Sharing bandwidth
- Two goals for bandwidth sharing
 - **Efficiency**: High utilization of link bandwidth
 - **Fairness**: Each flow gets equal share




Why AIMD?

- Every RTT, we can do
 - Multiplicative increase or decrease: $CWND \rightarrow a * CWND$
 - Additive increase or decrease: $CWND \rightarrow CWND + b$
- Four alternatives:
 - AIAD: gentle increase, gentle decrease
 - AIMD: gentle increase, drastic decrease
 - MIAD: drastic increase, gentle decrease
 - MIMD: drastic increase and decrease




Simple model of congestion control

- Two users
 - rates x_1 and x_2
- Congestion when $x_1 + x_2 > 1$
- Unused capacity when $x_1 + x_2 < 1$
- Fair when $x_1 = x_2$






Example






- Increase: $x + a_I$
- Decrease: $x - a_D$
- Does not converge to fairness






AIAD Sharing Dynamics





MIMD


- Increase: x^*b_I
- Decrease: x^*b_D
- Does not converge to fairness






AIMD

- Increase: $x + a_I$
- Decrease: $x * b_D$
- Converges to fairness






AIMD Sharing Dynamics



- Increase: x^*b_I
- Decrease: $x - a_D$
- Does not converge to fairness
- Does not converge to efficiency
- *"Analysis of the Increase and Decrease Algorithms for Congestion Avoidance in Computer Networks"*
-- Chiu and Jain





Summary

- Flow control ensures that the sender does not overflow the receiver
- Congestion control ensures that the sender does not overflow the network
 - Discover bandwidth
 - Adjust to conditions
 - Share bandwidth with others
- Slow Start
- AIMD



Homework

- Textbook Chapter 3: R14, P27, P32, P40, P45, P46, P50, P52