



南京大學

NANJING UNIVERSITY



Computer Networks

Wenzhong Li, Chen Tian

Nanjing University

Material with thanks to James F. Kurose, Mosharaf Chowdhury, and other colleagues.



Internet Applications

- WWW and HTTP
- Content Distribution Networks (CDNs)



WWW and HTTP



The Web: History

- World Wide Web (**WWW**): a distributed database of “pages” linked through Hypertext Transport Protocol (**HTTP**)
 - First HTTP implementation – 1990
 - **Tim Berners-Lee** at CERN
 - HTTP/0.9 – 1991
 - Simple GET command for the Web
 - HTTP/1.0 – 1992
 - Client/server information, simple caching

蒂姆·伯纳斯-李爵士
Sir Tim Berners-Lee



出生 1955年6月8日 (61歲)^[1]

✚ 英格兰伦敦

机构

万维网联盟
南安普敦大学
Plessey
麻省理工学院

知名于

发明万维网
麻省理工学院计算机科学及人工智能实
验室创办主席

2016 Turing Award



The Web: History (cont'd)

- World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)
 - HTTP/1.1 – 1996
 - Performance and security optimizations
 - HTTP/2 – 2015
 - Latency optimizations via request multiplexing over single TCP connection
 - Binary protocol instead of text
 - Server push



Web components

- Infrastructure:
 - Clients
 - Servers (DNS, CDN, Datacenters)
- Content:
 - URL: naming content
 - HTML: formatting content
- Protocol for exchanging information: HTTP




URL – Uniform Resource Locator

- A **unique identifier for an object** on WWW
- URL format
 - `<protocol>://<host>:<port>/<path>?query_string`
 - Protocol: method for transmission or interpretation of the object, e.g. http, ftp, Gopher
 - Host: DNS name or IP address of the host where object resides
 - Path: pathname of the file that contains the object
 - **Query_string**: name/value pairs sent to app on the server
- An example
 - `http://www.nju.edu.cn:8080/somedir/page.htm`



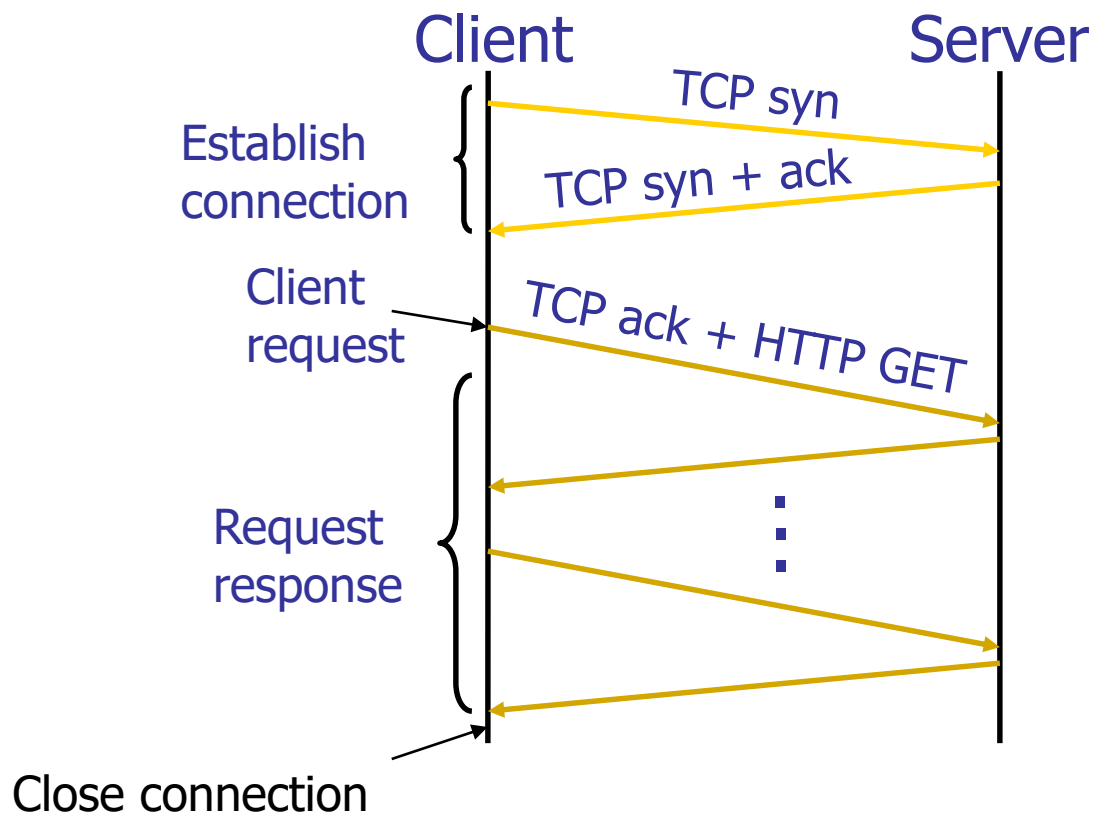
Hyper Text Transfer Protocol (HTTP)

- Client-server architecture
 - Server is “always on” and “well known”
 - Clients initiate contact to server
- Synchronous request/reply protocol
 - Runs over TCP, Port 80
- Stateless
- ASCII format
 - Before HTTP/2





Steps in HTTP request/response





Method types (HTTP 1.1)

- GET, HEAD
- POST
 - Send information (e.g., web forms)
- PUT
 - Uploads file in entity body to path specified in URL field
- DELETE
 - Deletes file specified in the URL field



Client-to-server communication

■ HTTP Request Message

- Request line: method, resource, and protocol version

request line — **GET /somedir/page.html HTTP/1.1**

header lines — **Host: www.someschool.edu**
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
(blank line)

*carriage return line feed
indicates end of message*



Server-to-client communication

■ HTTP Response Message

- Status line: protocol version, status code, status phrase
- Response headers: provide information
- Body: optional data

status line

(protocol, status code, status phrase)

header lines

data

e.g., requested HTML file

HTTP/1.1 200 OK

Connection close

Date: Thu, 06 Jan 2017 12:00:15 GMT

Server: Apache/1.3.0 (Unix)

Last-Modified: Mon, 22 Jun 2006 ...

Content-Length: 6821

Content-Type: text/html

(blank line)

data data data data data ...



HTTP is stateless

- Each request-response treated independently
 - Servers not required to retain state
- **Good**: Improves scalability on the server-side
 - Failure handling is easier
 - Can handle higher rate of requests
 - Order of requests doesn't matter
- **Bad**: Some applications need persistent state
 - Need to uniquely identify user or store temporary info
 - e.g., Shopping cart, user profiles, usage tracking, ...




Question

- How does a stateless protocol keep state?

State in a stateless protocol: Cookies

- Client-side state maintenance
 - Client stores small state on behalf of server
 - Client sends state in future requests to the server
- Can provide authentication





A Cookies Example

Client

Server



cookie file



ebay 8734
amazon 1678

one week later:



ebay 8734
amazon 1678

usual http request msg

usual http response
Set-cookie: 1678

usual http request msg
cookie: 1678

usual http response msg

usual http request msg
cookie: 1678

usual http response msg

Amazon server
creates ID
1678 for user


create
entry

cookie-
specific
action

cookie-
specific
action

access

access



backend
database



Application of Cookies

What cookies can bring

- Authorization
- Shopping carts
- Recommendations
- User session state (Web Email)

aside Cookies and privacy

- Cookies permit servers to learn a lot about user
- User may supply name and Email to servers
- Search engines may use cookies to obtain info across sites
- Hacked browser may do bad things with cookies



Web performance goals

- User
 - Fast downloads (not identical to low-latency communication!)
 - High availability
- Content provider
 - Happy users (hence, above)
 - Cost-effective infrastructure
- Network (secondary)
 - Avoid overload



Solutions?

■ User

- Fast downloads (not communication!) latency
- High availability

Improve networking protocols including HTTP, TCP, etc.

■ Content provider

- Happy users (hence, above)
- Cost-effective infrastructure

■ Network (secondary)

- Avoid overload



Solutions?

■ User

- Fast downloads (not communication!)

- High availability

■ Content provider

- Happy users (hence, above)
- Cost-effective infrastructure

■ Network (secondary)

- Avoid overload

Improve networking protocols including HTTP, TCP, etc.

Latency

Caching and replication



Solutions?

■ User

- Fast downloads (not communication!) latency
- High availability

Improve networking protocols including HTTP, TCP, etc.

■ Content provider

- Happy users (hence, above)
- Cost-effective infrastructure

Caching and replication

■ Network (secondary)

- Avoid overload

Exploit economies of scale; e.g., webhosting, CDNs, datacenters




HTTP performance

- Most Web pages have multiple objects
 - e.g., HTML file and a bunch of embedded images
- How do you retrieve those objects (naively)?
 - One item at a time
- New TCP connection per (small) object!



Object request response time

- RTT (round-trip time)
 - Time for a small packet to travel from client to server and back
- Response time
 - 1 RTT for TCP setup
 - 1 RTT for HTTP request and first few bytes
 - Transmission time
 - **Total** = 2RTT + Transmission Time





Non-persistent connections


- Default in HTTP/1.0
- $2RTT + \triangle$ for each object in the HTML file!
 - One more $2RTT + \triangle$ for the HTML file itself
- Doing the same thing over and over again
 - Inefficient



Concurrent requests and responses

- Use multiple connections in parallel
- Does not necessarily maintain order of responses

- Client = 😊
- Content provider = 😊
- Network = 😞 Why?





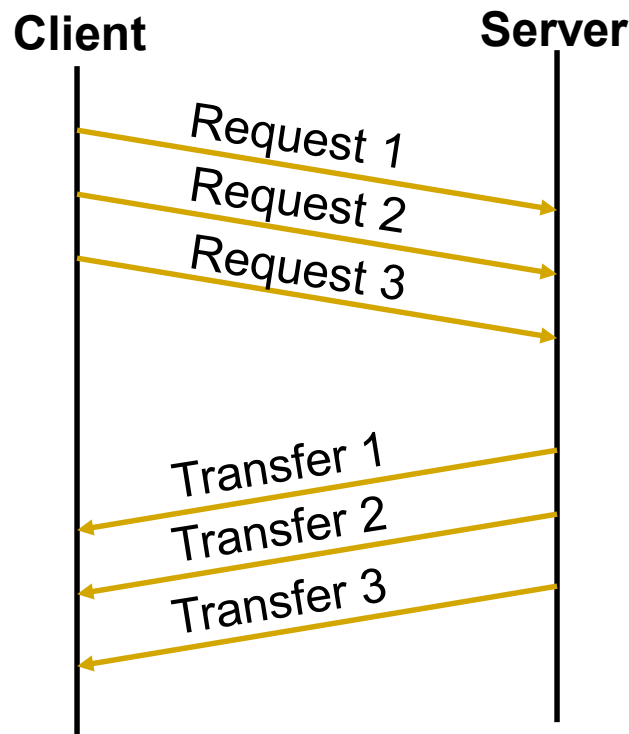
Persistent connections

- Maintain TCP connection across multiple requests
 - Including transfers subsequent to current page
 - Client or server can tear down connection
- Advantages
 - Avoid overhead of connection set-up and tear-down
 - Allow underlying layers (e.g., TCP) to learn about RTT and bandwidth characteristics
- Default in HTTP/1.1



Pipelined requests & responses

- Batch requests and responses to reduce the number of packets
- Multiple requests can be contained in one TCP segment





Scorecard: Getting n small objects

- Time dominated by latency
- One-at-a-time: $\sim 2n$ RTT
- m concurrent: $\sim 2\lceil n/m \rceil$ RTT
- Persistent: $\sim (n+1)$ RTT
- Pipelined: ~ 2 RTT
- Pipelined/Persistent: ~ 2 RTT first time, RTT later



Scorecard: Getting n large objects each of size F

- Time dominated by bandwidth
- One-at-a-time: $\sim nF/B$
- m concurrent: $\sim [n/m] F/B$
 - Assuming shared with large population of users and each TCP connection gets the same bandwidth
- Pipelined and/or persistent: $\sim nF/B$
 - The only thing that helps is getting more bandwidth



Caching

- Why does caching work?
 - Exploits locality of reference
- How well does caching work?
 - Very well, up to a limit
 - Large overlap in content
 - But many unique requests
 - A universal story!
 - Effectiveness of caching grows logarithmically with size



Caching: How

- Modifier to GET requests:
 - **If-modified-since** – returns “not modified” if resource not modified since specified time

GET /somedir/page.html HTTP/1.1

Host: www.someschool.edu

User-agent: Mozilla/4.0

If-modified-since: Wed, 18 Jan 2017 10:25:50 GMT
(blank line)



Caching: How

- Modifier to GET requests:
 - **If-modified-since** – returns “not modified” if resource not modified since specified time
- Client specifies “**if-modified-since**” time in request
- Server compares this against “last modified” time of resource
- Server returns “Not Modified” if resource has not changed
- or a “OK” with the latest version otherwise



Caching: How

- Modifier to GET requests:
 - **If-modified-since** – returns “not modified” if resource not modified since specified time
- Response header:
 - **Expires** – how long it’s safe to cache the resource
 - **No-cache** – ignore all caches; always get resource directly from server




Caching: Where?

- Options
 - Client (browser)
 - Forward proxies
 - Reverse proxies
 - Content Distribution Network



Caching: Where?


- Many clients transfer same information
 - Generate unnecessary server and network load
 - Clients experience unnecessary latency





Caching with Reverse Proxies


- Cache documents close to server
 - Decrease server load
 - By content provider





Caching with Forward Proxies

- Cache documents close to clients
 - Reduce network traffic and decrease latency
 - By ISPs or enterprises





- HTTP/1.1
 - Text-based protocol
 - Being replaced by binary HTTP/2 protocol
- Many ways to improve performance
 - Pipelining and batching
 - Caching in proxies and CDNs
 - Datacenters



Content Distribution Networks (CDNs)




Content Distribution Networks (CDNs)

■ Challenge

- Stream large files (e.g. video) from single origin server in real time
- Protect origin server from DDOS attacks

■ Solution

- Replicate content at **hundreds of servers** throughout Internet
- **CDN distribution node** coordinate the content distribution
- Placing content **close to user**





Content Replication

- Content provider (origin server) is **CDN customer**
- CDN replicates customers' content in CDN servers
- When provider updates content, CDN updates its servers
- Use **authoritative DNS server** to redirect requests



Supporting Techniques

■ DNS

- One name maps onto many addresses

■ Routing

- Content-based routing (to nearest CDN server)


■ URL Rewriting

- Replaces "http://www.sina.com/sports/tennis.mov" with "http://www.cdn.com/www.sina.com/sports/tennis.mov"

■ Redirection strategy

- Load balancing, network delay, cache/content locality

CDN Operation



- 1' URL rewriting – get authoritative server
1. Get near CDN server IP address
2. Warm up CDN cache
3. Retrieve pages/media from CDN Server



Redirection

- CDN creates a “**map**”, indicating distances from leaf ISPs and CDN servers
- When query arrives at **authoritative DNS server**
 - Server determines ISP from which query originates
 - Uses “map” to determine best CDN server
- CDN servers create an **application-layer overlay network**



Summary of Web App

- Conceptual, implementation aspects of network application protocols
 - Client-Server vs. Peer-to-Peer
 - Data presentation formatting
- Examining popular application-level protocols
 - DNS, SNMP / MIB
 - HTTP, FTP, SMTP / POP3 / MIME
 - Content distribution networks (CDNs)



Summary

- Internet Applications
 - WWW and HTTP
 - Content Distribution Networks (CDNs)